

# CS4223

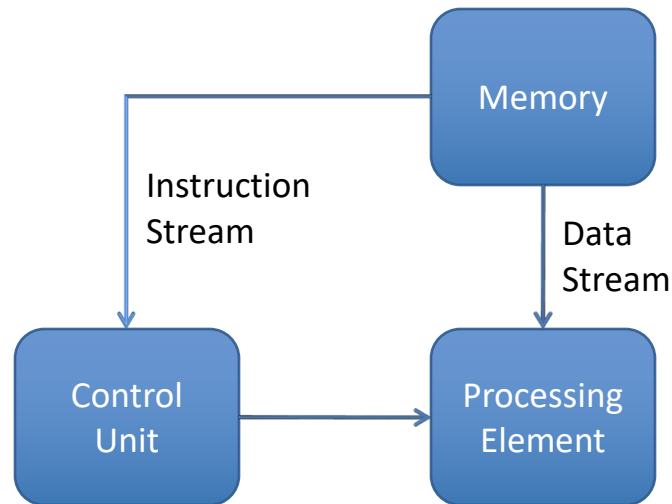
## Many- and multi-core for TLP

Trevor E. Carlson  
National University of Singapore  
tcarlson@comp.nus.edu.sg  
(Slides from Tulika Mitra)

# Flynn's Taxonomy of Parallel Computers

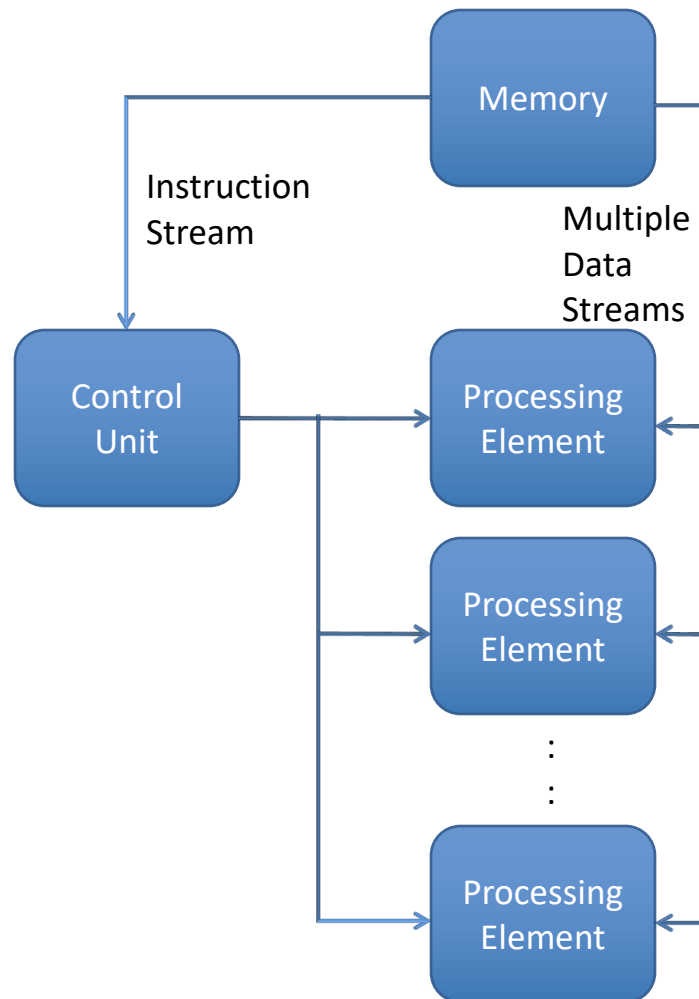
		Number of Data Streams	
		Single	Multiple
Number of instruction streams	Single	SISD	SIMD
	Multiple	MISD	MIMD

# SISD: Single Instruction Single Data



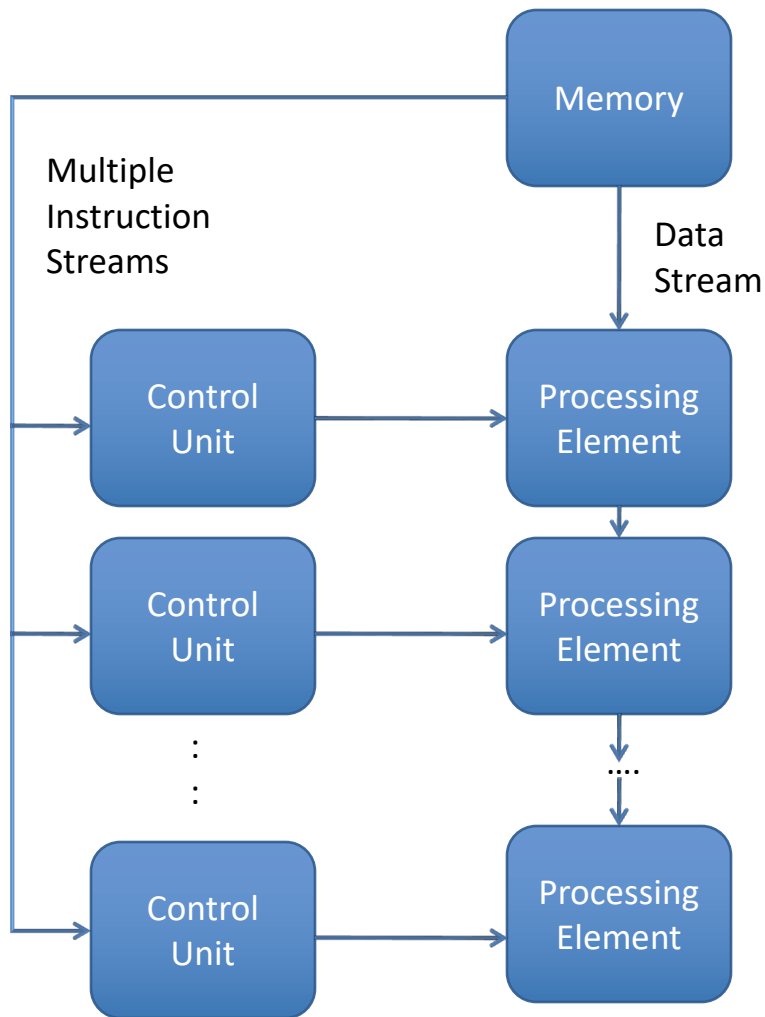
- SISD is not considered as a parallel architecture
- SISD exploits parallelism at the instruction level
- Pipelined, Superscalar, and VLIW architectures are examples of SISD architecture

# SIMD: Single Instruction Multiple Data



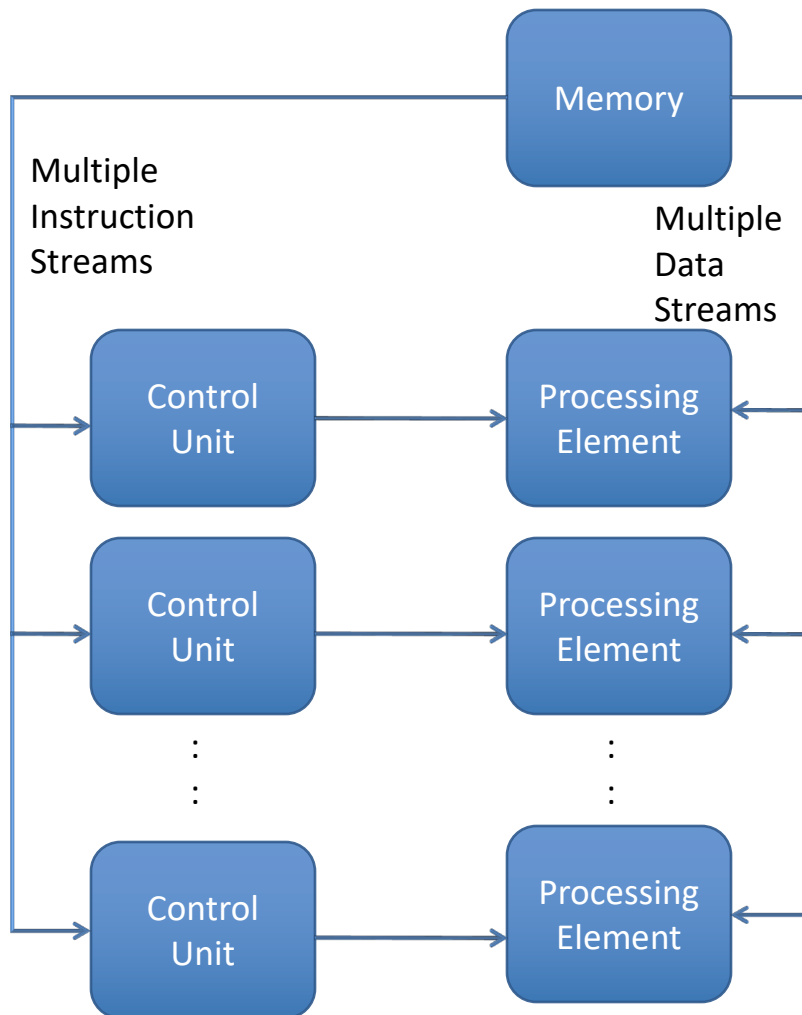
- A single instruction operates on multiple data to exploit data parallelism
- Vector processors and GPUs are excellent examples of SIMD architecture
- More efficient in terms of instruction count and loop control overhead

# MISD: Multiple Instruction Single Data



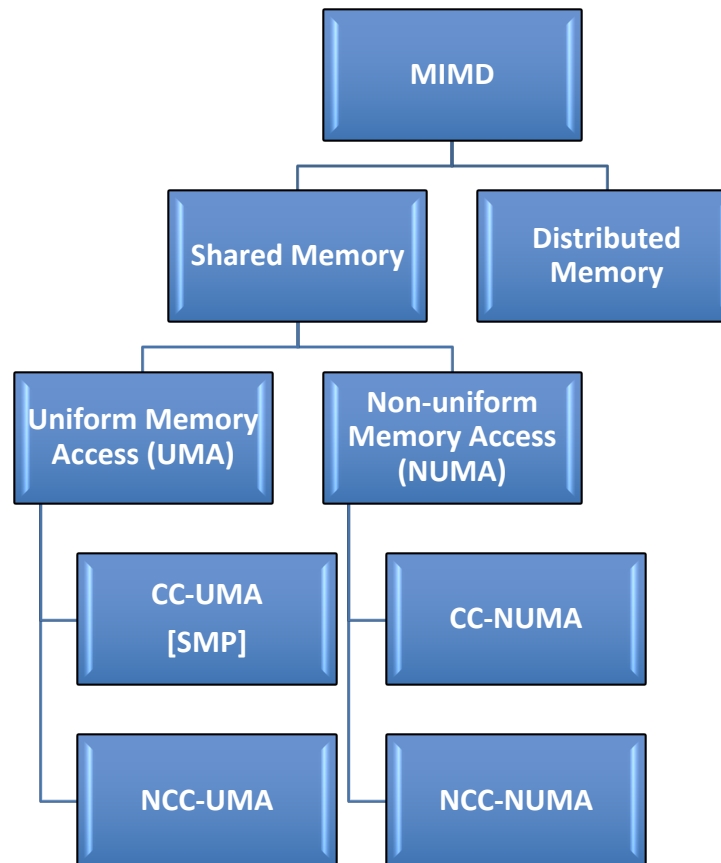
- Multiple processing elements execute from different instruction streams and data is passed from one processing element to the next
- Example: Systolic array such as CMU iWrap
- Data passing restriction is quite severe --- hard to generalize

# MIMD: Multiple Instruction Multiple Data




- Most flexible architecture
- Used in most parallel computers today

# MIMD Classification



# Distributed Memory MIMD

- Processing elements (PE) work independently
- Interaction is via message passing
- PEs cannot access the memory of another PE directly





# Distributed Memory MIMD (contd.)


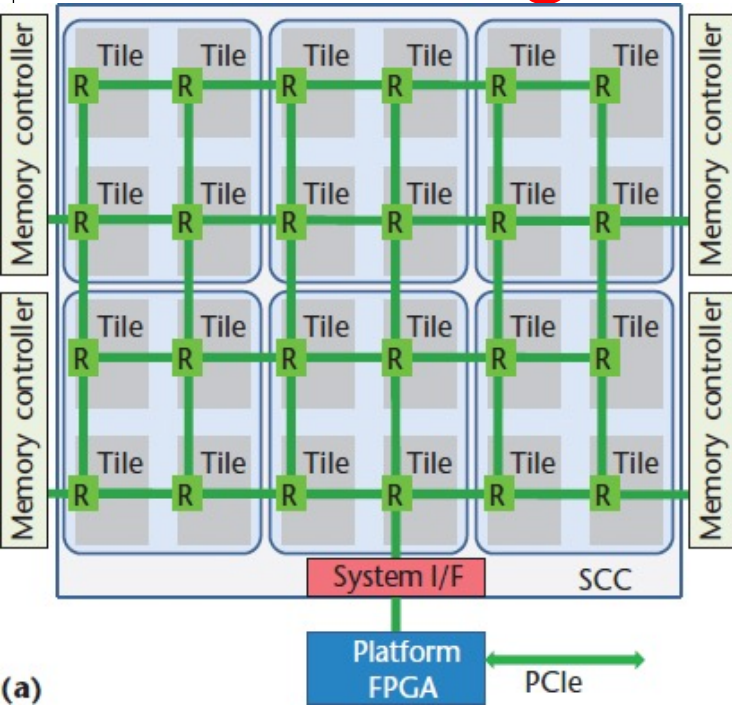
- Advantages

- Processors mostly work with local memory
  - less contention and highly scalable
- As communication is via message passing, sophisticated synchronization techniques are not required

- Disadvantages


- Load balancing (partition of code and data) responsibility is on the programmer
- Message-passing-based synchronization can lead to deadlocks; programmer's responsibility
- Performance overhead due to physical data copy

# Intel Single Chip Cloud Computer (SCC)



# Shared Memory MIMD

- Multiple processors with a **logically** shared memory
- The memory modules define a global address space that is shared among the processors
- Any processor can access any memory module
- Communication through memory loads and stores
- Logically shared memory can be physically distributed




# Shared Memory MIMD (contd.)

- Advantages:
  - No need to partition code or data
  - No need to physically move data among processors
    - communication is efficient
- Disadvantages:
  - Special synchronization constructs are required
  - Lack of scalability due to contention

# Shared Memory MIMD Classification


- Processor has cache ( CC / NCC )
  - Same shared variable may exist in multiple caches
  - Hardware ensures correctness via cache coherence protocol
- Processor to memory delay ( UMA / NUMA )
  - All accesses to memory encounter same delay or not [ uniform versus non-uniform memory access]



# Cache Coherence

- Caches reduce data access time and bandwidth requirement on the interconnect
- But private caches create problem
  - Copies of a variable in multiple caches
  - A write by one processor may not become visible to others
    - They may keep accessing stale data in their caches
  - Cache coherence problem

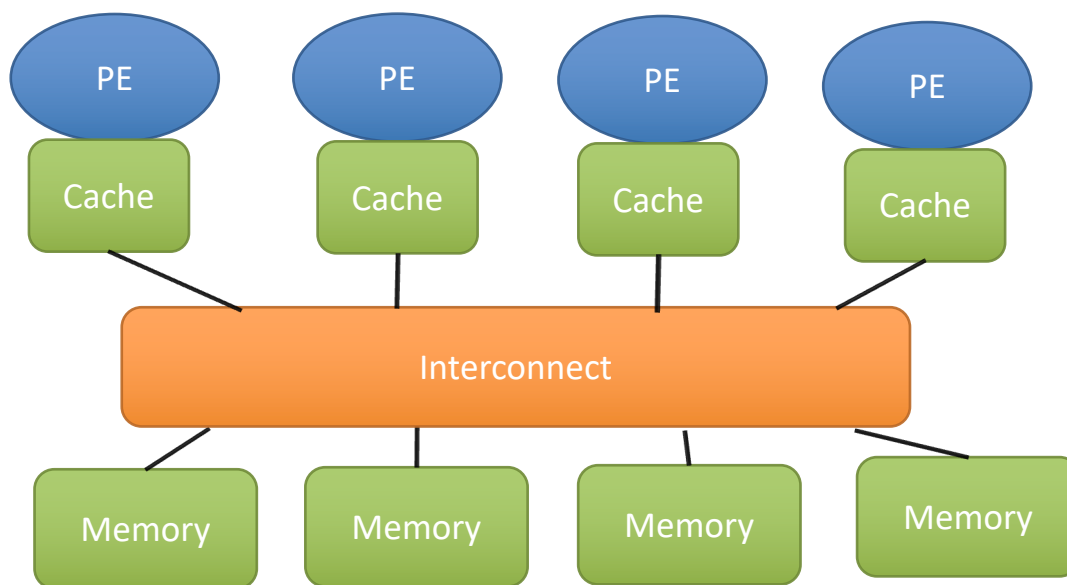
# Cache Coherence Problem



- Processors see different values of  $u$  after event 3
- With write-back caches, main memory will have the stale value till the entry is replaced


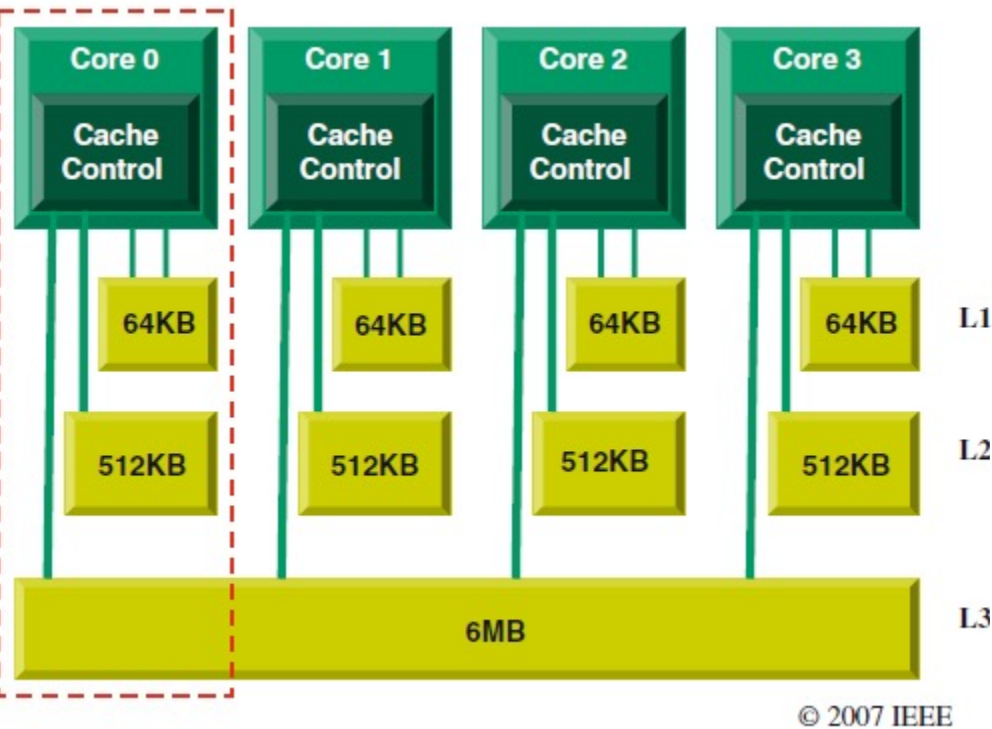
# CC-UMA or SMP

- Cache-Coherent Uniform Memory Access
- Also known as Symmetric Multiprocessor (SMP)

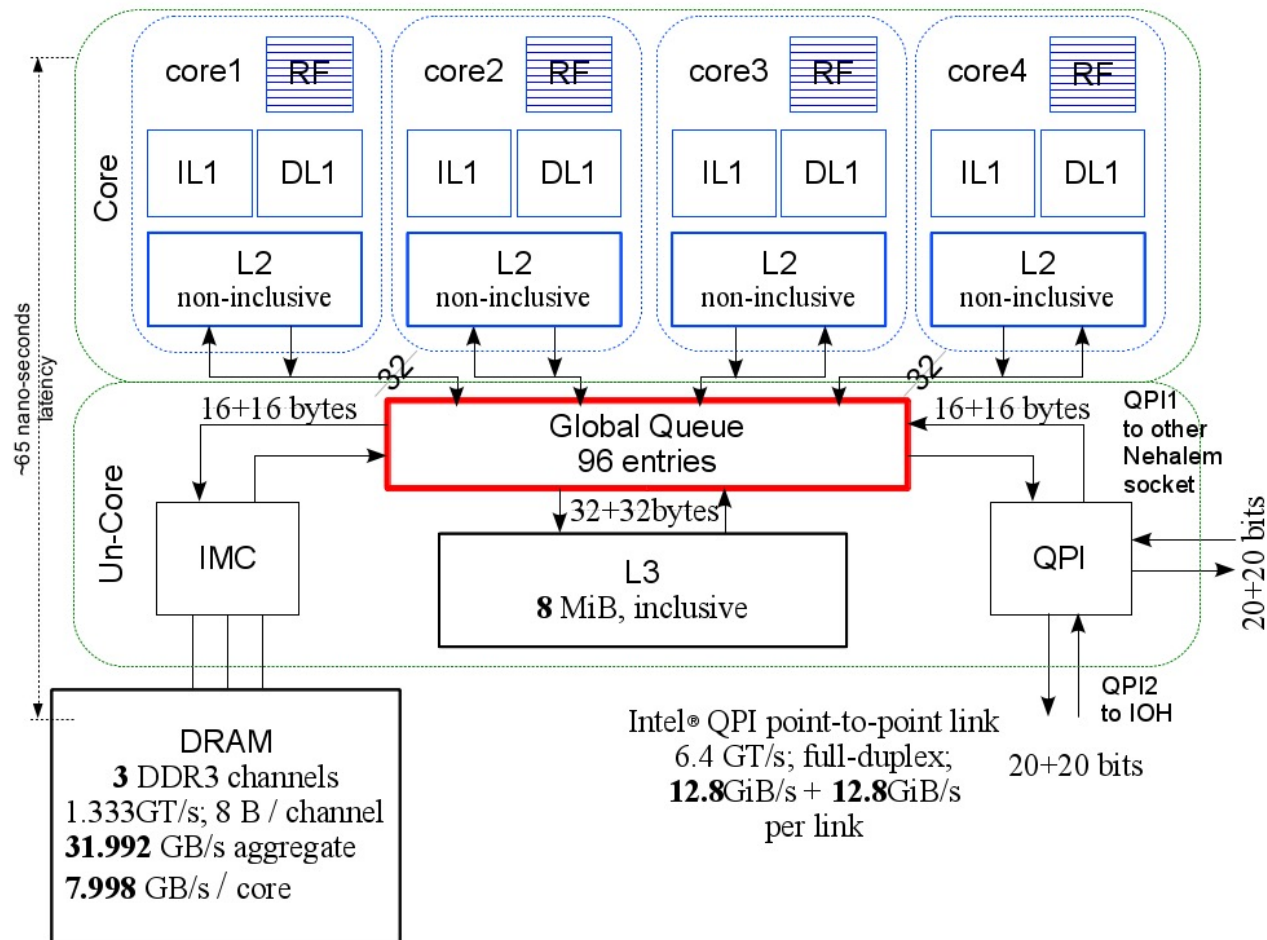




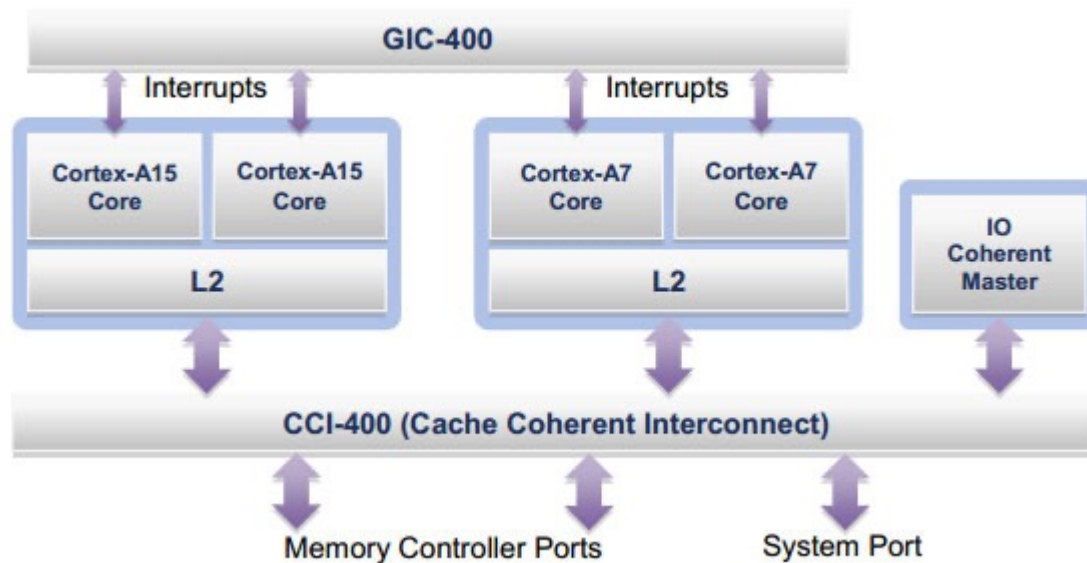
# Contemporary CC-UMA: AMD Opteron



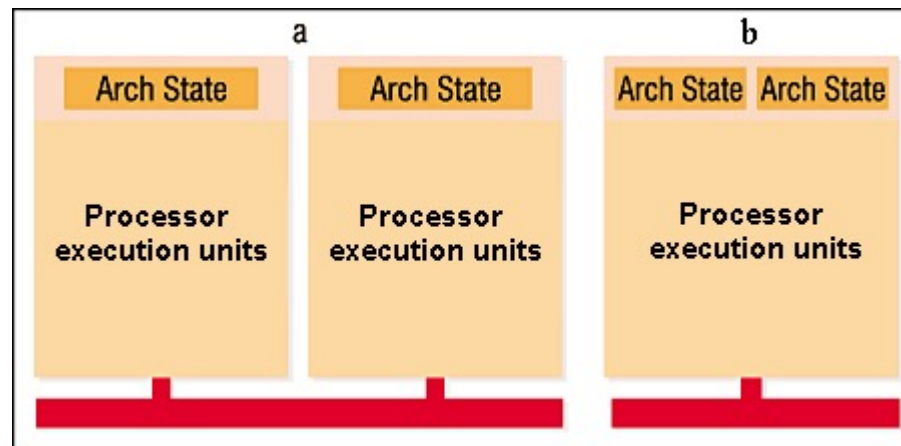
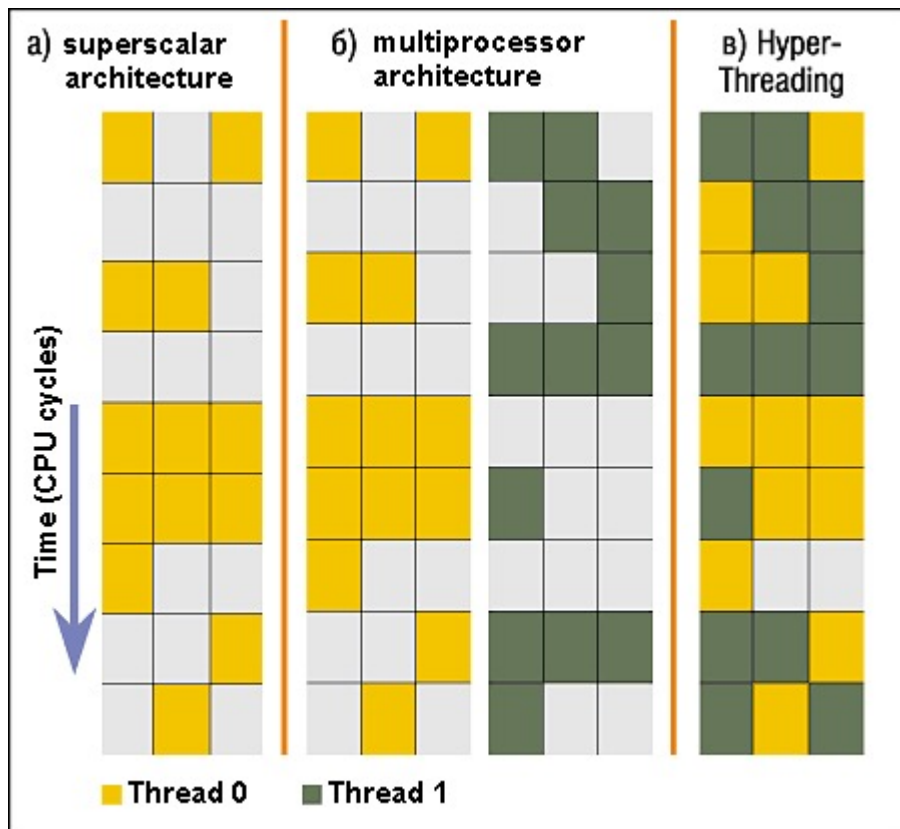
# Contemporary CC-UMA: Intel Nehalem



# Mobile CC-UMA: ARM big.LITTLE




# Simultaneous Multi-threading SMT (a.k.a. Hyper-threading)



# Hyper-threading demo


<https://www.youtube.com/watch?v=ptnotrfmA1E>

# CC-UMA: Scaling to Many Core Tiler




# CC-NUMA

- Cache Coherent Non Uniform Memory Access
- Each node has cache memory to reduce contention

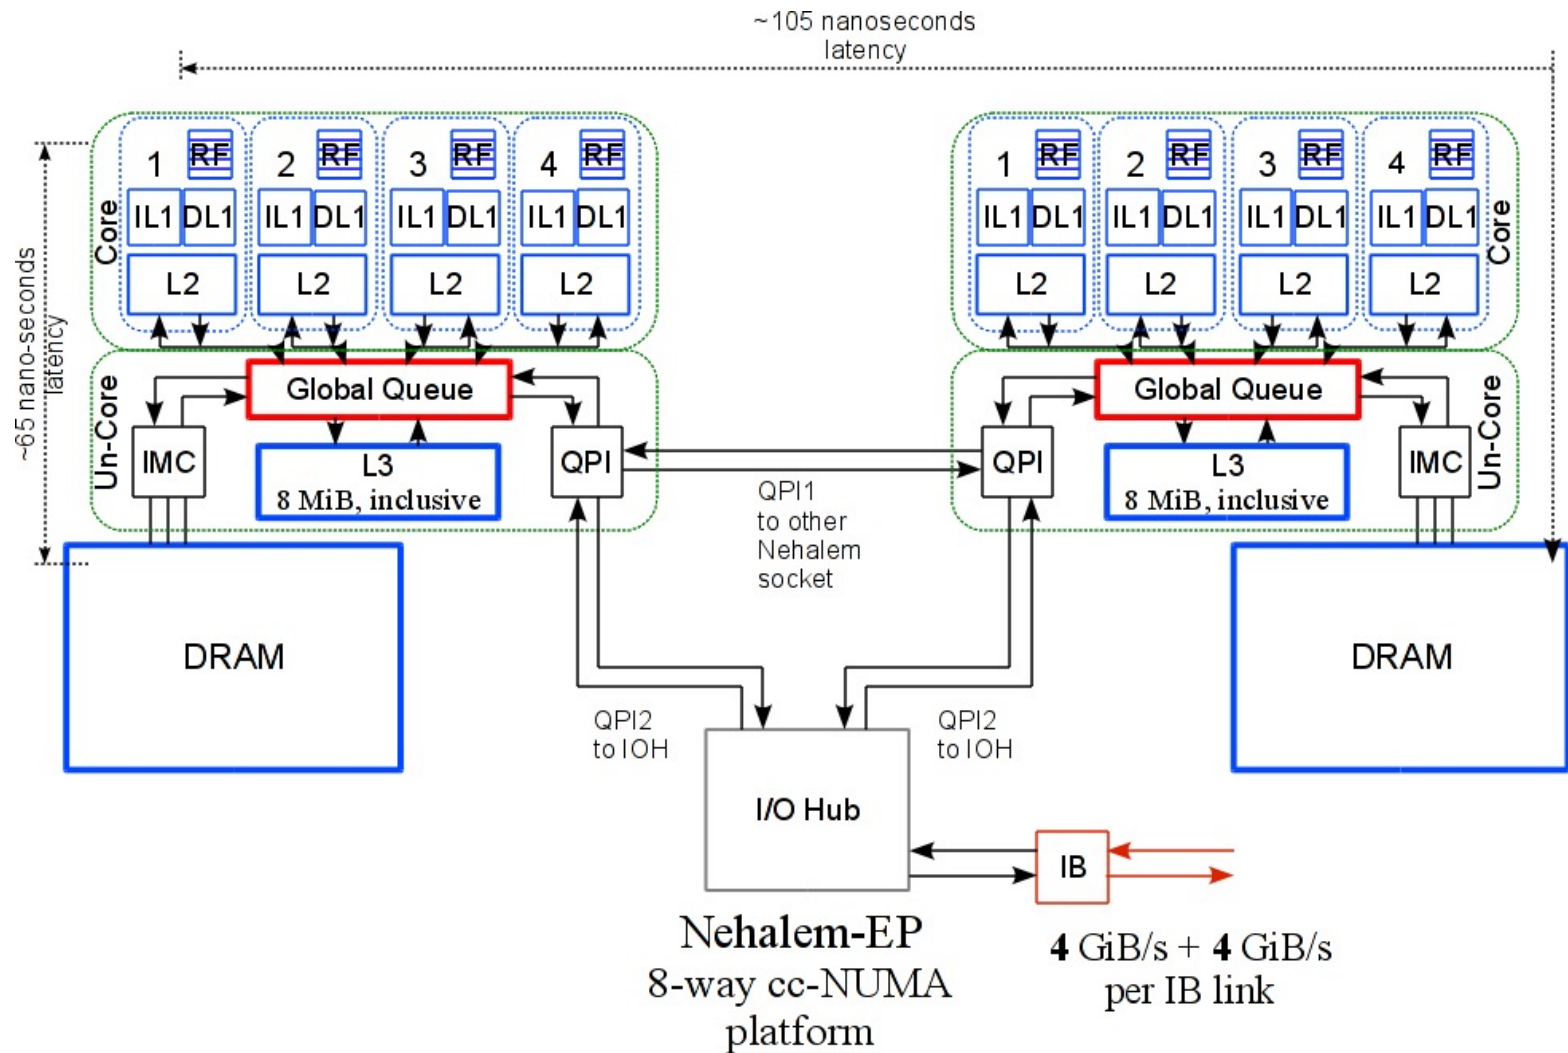


# cc-NUMA: AMD 6100 Opteron package






# CC-NUMA: Nehalem Cluster




# NCC-UMA

- Non cache-coherent uniform memory access
- Too much memory contention
- Does not make sense



# NCC-NUMA

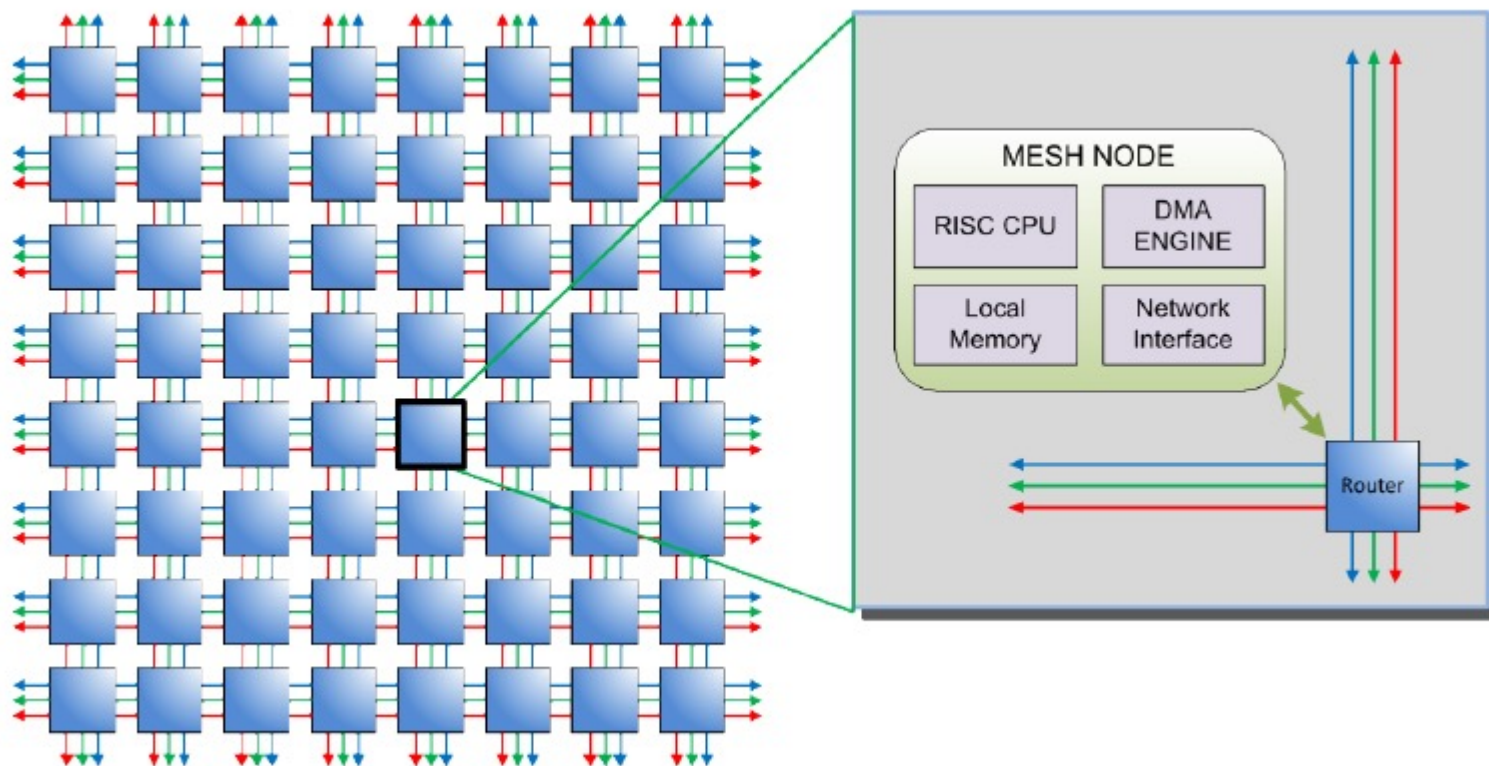
- Non Cache-Coherent Non-Uniform Memory Access
- Memory is logically shared, physically distributed
- Requires careful load balancing and data distribution just as distributed memory MIMD; programming is still based on shared memory



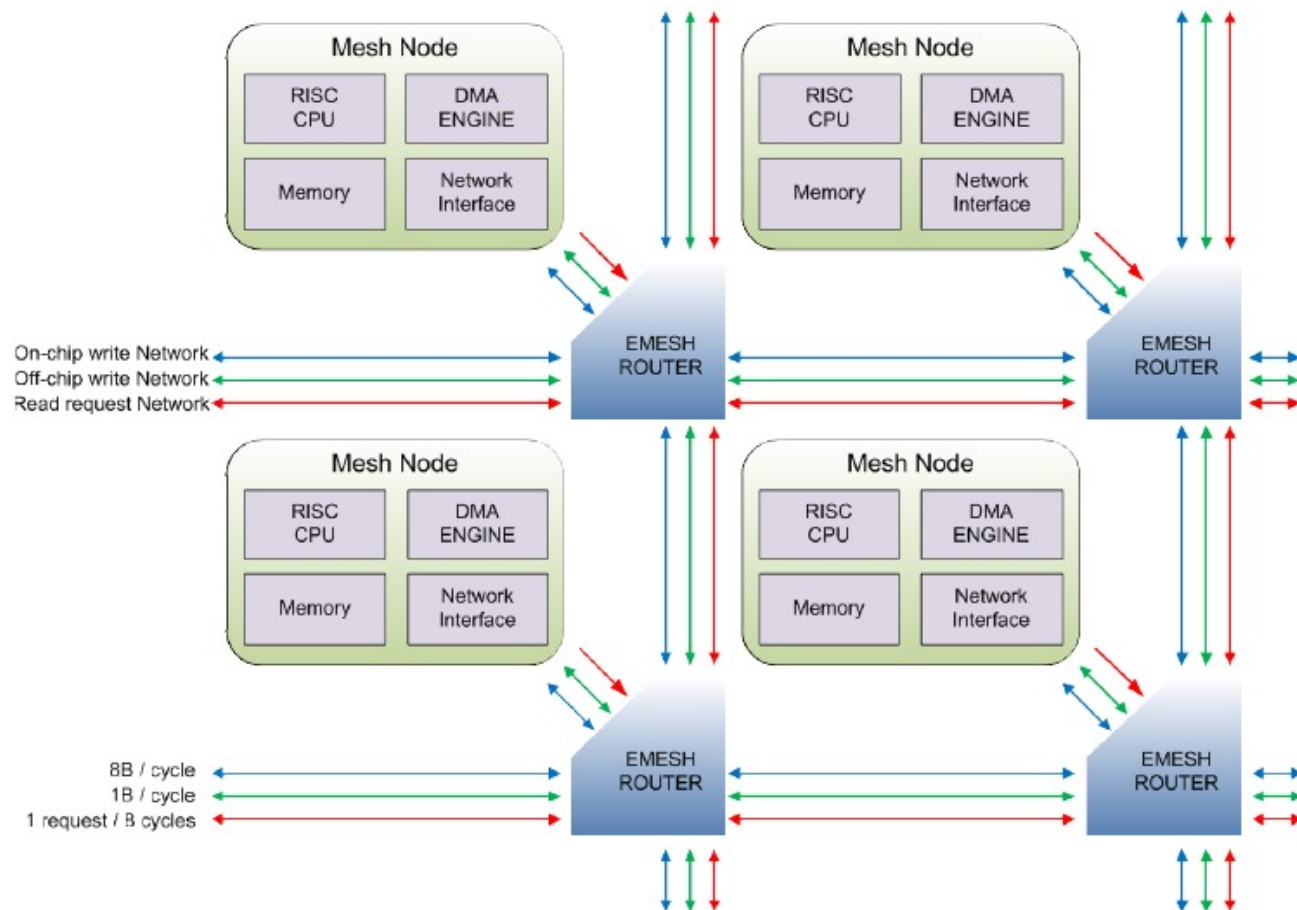
# Distributed shared memory MIMD vs. Distributed memory MIMD

- Physical construction is the same
- In distributed shared memory MIMD (NUMA), local memories are components of global address space and any processor can access the local memory of any other processor directly
- In distributed memory MIMD, local memories have separate address spaces, and direct access to the local memory of a remote processor is prohibited

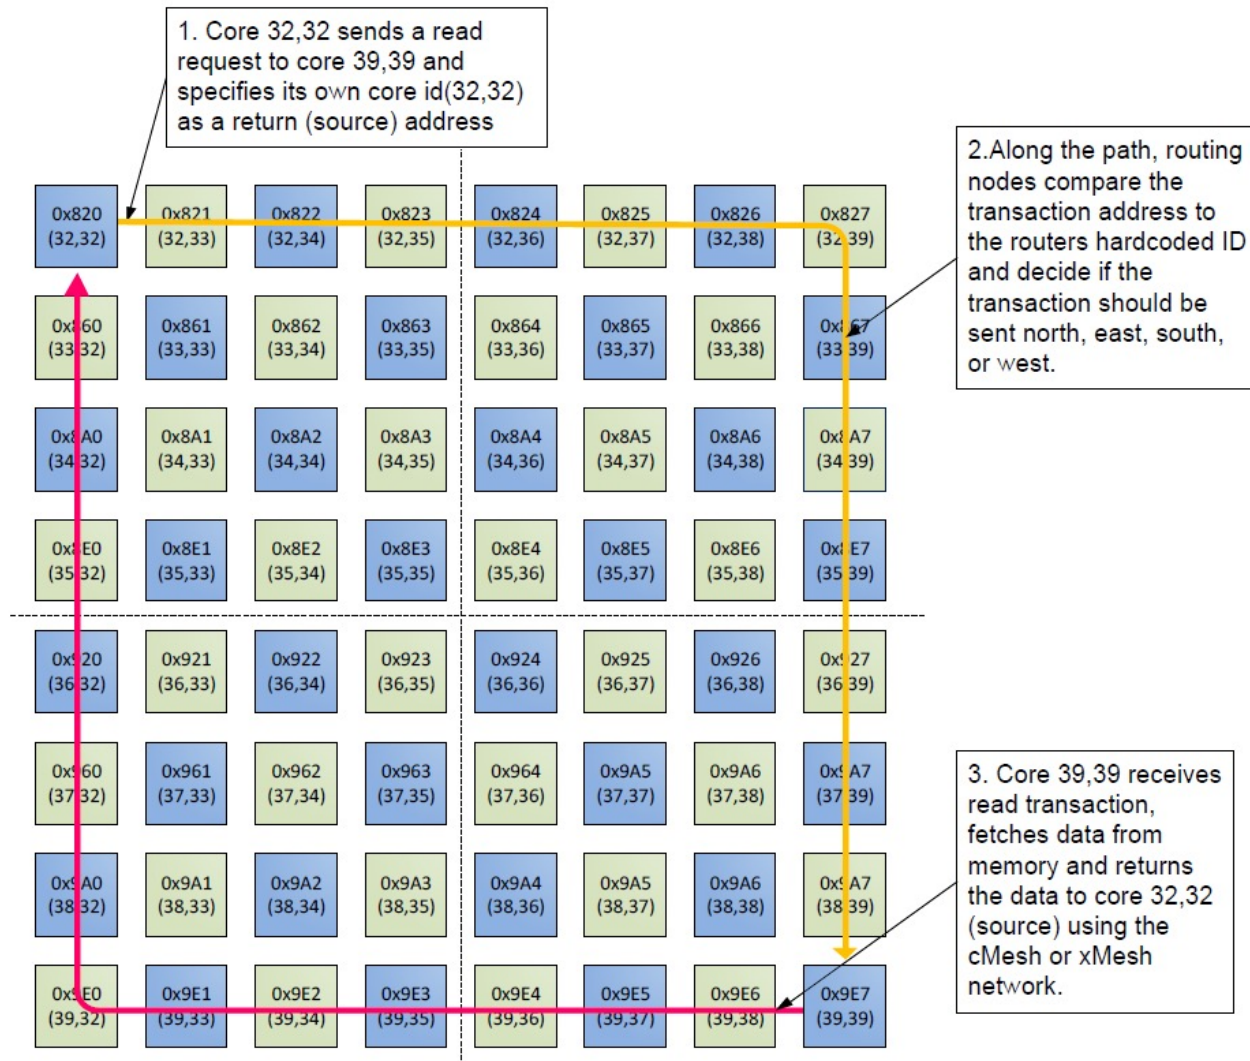
# There's Something about Many: Parallel Epiphany Architecture



# Network-on-chip



# Routing



# Making of Intel Core i7

<http://www.youtube.com/watch?v=tKX8bdHWgu8>