

CS4223

Synchronization

Trevor E. Carlson
National University of Singapore
tcarlson@comp.nus.edu.sg
(Slides from Tulika Mitra)

Learning Objectives

- Why do we need synchronization?
- Building block: Atomic read-modify-write instructions
- Lock implementations with increasing complexity:
 - Test & Set lock
 - Test and Test & Set Lock
 - Locks based on LL-SC primitive
- Fair lock implementations:
 - Ticket lock
 - Array-based queuing lock
- Barrier implementation

Synchronization

- A parallel computer is a collection of processing elements that **cooperate** and communicate to solve large problems fast
- Synchronization for co-operation
 - Mutual exclusion (**locks**)
 - Event synchronization
 - Point-to-point
 - Global (**barrier**)

Shared memory code

```
begin parallel // spawn a child thread
private int start_iter, end_iter, i;
shared int local_iter = 4;
shared double sum = 0.0, a[], b[], c[];
shared lock_type mylock;

start_iter = gettid() * local_iter;
end_iter = start_iter + local_iter;
for (i=start_iter; i<end_iter; i++)
    a[i] = b[i] + c[i];
barrier;

for (i=start_iter; i<end_iter; i++)
    if (a[i] > 0) {
        lock (mylock);
        sum = sum + a[i];
        unlock (mylock);
    }
barrier;
end parallel // kill the child thread
print sum;
```

First attempt at simple software lock

```
lock:ld register, location /* copy location to register */  
     cmp location, #0      /* compare with 0 */  
     bnz lock              /* if not 0, try again */  
     st location, #1       /* store 1 to mark it locked */  
     ret                   /* return control to caller */  
  
unlock:st    location, #0 /* write 0 to location */  
         ret              /* return control to caller */
```

- Problem: lock needs atomicity in its implementation
 - Read (test) and write (set) of lock variable by a process not atomic
- Solution: atomic read-modify-write instruction
 - Atomically test value of location and set it to another value, return success or failure

Atomic Read-Modify-Write Instruction

- Specifies a location and register. In atomic operation:
 - Value in location read into a register
 - Another value stored into location
- Many variants
 - Varying degrees of flexibility in second part
- Simple example: **test & set**
 - Value in location read into a specified register
 - Constant 1 stored into location
 - Successful if value loaded into register is 0
- Can be used to build locks

How is atomicity ensured?

- Rely on cache coherence
- Obtain exclusive ownership of memory location M
- Do not allow the block to be stolen during atomic instruction execution
 - For the duration of the atomic instruction execution, lock the bus
 - Alternatively, the cache controller of the processor executing atomic instruction defer responding to all requests to the block until the atomic instruction completes execution

Other Read-Modify-Write Primitives

- **Exchange Rx, M:**
 - Atomically exchange the value in M with value in Rx
- **Fetch & op M:**
 - Read value stored in memory location M
 - Perform op (increment, decrement, addition, subtraction) to it
 - Store the new value to memory location M
- **Compare & swap Rx, Ry, M**
 - Three operands: location, register to compare , register to swap
 - Compare value in M with value in register Rx
 - If they match, write value in Ry to M
 - Copy the value in Rx to Ry

Simple Test & Set Lock

```
lock:  t&s register, location
       bnz lock           /* if not 0, try again */
       ret                /* return control to caller */

unlock:st location, #0    /* write 0 to location */
       ret                /* return control to caller */
```

Simple T&S Performance

- Un-contended lock-acquisition latency low
 - One atomic instruction + branch
- Traffic requirement: high
 - Each lock acquisition attempt causes invalidation of all cached copies regardless of whether acquisition is successful or not

Test & Set with backoff

- Reduce frequency of issuing t&s while waiting
 - test & set lock with backoff
 - Don't back off too much or will be in backed off state when lock becomes free
 - Exponential backoff works quite well empirically:
 $i^{\text{th}} \text{ time} = k * c^i$

Test and Test & Set Lock

- Busy-wait with read operations rather than t&s
- Keep testing with ordinary load
 - cached lock variable invalidated when release occurs
- When value changes to 0, try to obtain lock with t&s
 - only one processor will succeed
 - others will fail and start testing again

Test and Test & Set lock

```
lock: ld register, location  /* copy location to register */  
      cmp register, #0      /* compare with 0 */  
      bnz lock              /* if not 0, try again */  
      t&s register, location  
      bnz lock  
      ret                  /* return control to caller */  
  
unlock: st location, #0     /* write 0 to location */  
        ret                /* return control to caller */
```

Test and Test & Set lock drawbacks

- All processor still rush out to read miss and t&s on release of the lock
- Poor fairness; can cause starvation

Load-Locked (Linked) Store-Conditional

- Goals:
 - Test with reads
 - Failed read-modify-write attempts do not generate invalidations
 - single primitive can implement range of r-m-w operations

LL-SC

- LL reads variable into register
- Follow with arbitrary instructions to manipulate its value
- SC tries to store back to location if and only if no one else has written to the variable since this processor's LL
 - If SC succeeds, means all three steps happened atomically
 - If fails, doesn't write or generate invalidations (retry LL)
 - Success indicated by condition codes

Implementation of LL/SC

- LL loads a block into register
- Records address in a special register **linked register**
- SC succeeds only if the address matches the address stored in linked register
- **Linked register is cleared on**
 - **An invalidation to the linked register address**
 - Context switching
- When SC fails, the store is canceled
 - Does not generate bus transaction

Simple Lock with LL-SC

```
lock:      ll      reg1, location    /* LL location to reg1 */
           bnz     reg1, lock
           sc      location, reg2    /* SC reg2 into location */
           beqz    lock              /* if failed, start again */
           ret
unlock:    st      location, #0      /* write 0 to location */
           ret
```

Finer details

- More fancy atomic ops by changing what is between LL-SC
 - Keep it small so SC likely to succeed
 - Do not include instructions that would need to be undone (e.g., stores)
- SC can fail (without putting transaction on bus) if:
 - Detects intervening write even before trying to get bus
 - Tries to get bus but another processor's SC gets bus first
- LL, SC are not lock, unlock respectively
 - Only guarantee no conflicting write to lock variable between them
 - Can be used directly to implement simple operations on shared variables

LL-SC Lock drawbacks

- No invalidation on failure, but read misses by all waiters after both release and successful SC by winner
- Does not reduce traffic to minimum and not a fair lock
- How to improve?
 - Fair lock
 - Only one processor to have read miss upon release and success

Drawbacks of earlier solutions

- Locks are not fair
- Lot of traffic on successful acquire and release due to cache miss effect

How to introduce fairness?

- Maintain the order in which the threads attempt to acquire lock
- On release, ensure that the first thread that asked for the lock acquires it

next-ticket

now-serving

Ticket Lock

- Works like waiting line at deli or bank
- Two counters per lock: `next_ticket`, `now_serving`
- Acquire:
fetch&inc `next_ticket`; wait for `now_serving` to equal it
- Release: increment `now-serving`
- Atomic operation when arrive at lock; not for release
- FIFO order: fair
- Unlike LL-SC lock, no invalidation when succeeds
- Still read misses at release; all spin on same variable

Ticket lock implementation


```
ticketLock_init(int *next_ticket, int *now_serving)
{
    *now_serving = *next_ticket = 0;
}
```

```
ticketLock_acquire(int *next_ticket, int *now_serving)
{
    my_ticket = fetch&inc(next_ticket);
    while (*now_serving != my_ticket) {};
}
```

```
ticketLock_release(int *now_serving)
{
    *now_serving++;
}
```


How to avoid read miss on release?

- Make each process wait on a different variable rather than a single variable
- Create array of N bits when N is no of threads



Array-based Queuing Locks

- Waiting processes poll on different locations in array
- Acquire
 - fetch&inc to obtain next array element to spin on
 - Ensure array elements are in different memory blocks
- Release
 - set next location in array, waking up process spinning on it
- $O(1)$ traffic per acquire with coherent caches
- FIFO ordering as in ticket lock
- $O(p)$ space per lock

Array-based Queuing Lock Implementation

```
init(int *next_ticket, int *can_serve){
    *next_ticket = 0;
    for (i=1; i<MAXSIZE; i++) can_serve[i] = 0;
    can_serve[0] = 1;
}

acquire(int *next_ticket, int *can_serve){
    my_ticket = fetch&inc(next_ticket);
    while (can_serve[my_ticket] != 1){};
}

Release(int *can_serve){
    can_serve[my_ticket + 1] = 1;
    can_serve[my_ticket] = 0;
}
```

Point to Point Event Synchronization

- Software methods:
 - Interrupts
 - Busy-waiting: use ordinary variables as flags

P1

`a = f(x); /* set a */`

`flag = 1;`

P2

`while (flag == 0) do nothing;`

`b = g(a); /* use a */`


Shared memory code

```
begin parallel // spawn a child thread
private int start_iter, end_iter, i;
shared int local_iter = 4;
shared double sum = 0.0, a[], b[], c[];
shared lock_type mylock;

start_iter = getid() * local_iter;
end_iter = start_iter + local_iter;
for (i=start_iter; i<end_iter; i++)
    a[i] = b[i] + c[i];
barrier;


for (i=start_iter; i<end_iter; i++)
    if (a[i] > 0) {
        lock (mylock);
        sum = sum + a[i];
        unlock (mylock);
    }
barrier;
end parallel // kill the child thread
print sum;
```

The idea of barrier



- Keep count of how many threads have reached barrier
- When count = N, threads can go past barrier

Simple centralized barrier




- First thread sets `flag = 0`
- Keep count of how many threads have reached barrier
- When `count = N`, set `flag = 1` releasing barrier

A simple centralized barrier

```
struct bar_type {
    int counter; struct lock_type lock; int flag = 0;
} bar_name;

BARRIER (bar_name, p) {
    LOCK(bar_name.lock);
    if (bar_name.counter == 0)
        bar_name.flag = 0;           /* reset flag if first to reach */
    mycount = bar_name.counter++;     /* mycount is private */
    UNLOCK(bar_name.lock);
    if (mycount == p) {               /* last to arrive */
        bar_name.counter = 0;         /* reset for next barrier */
        bar_name.flag = 1;           /* release waiters */
    }
    else while (bar_name.flag == 0) {}; /* busy wait for release */
}
```



Simple centralized barrier problem



A Working Centralized Barrier

- Consecutively entering the same barrier doesn't work
 - Must prevent process from entering until all have left previous instance
- Sense reversal: wait for flag to take different value consecutive times
 - Toggle this value only when all processes reach

Sense reversal barrier



Sense-reversal Barrier

```
BARRIER (bar_name, p) {  
    local_sense = !(local_sense); /* toggle private sense variable */  
    LOCK(bar_name.lock);  
    bar_name.counter++;  
    if (bar_name.counter == p){  
        UNLOCK(bar_name.lock);  
        bar_name.counter = 0;  
        bar_name.flag = local_sense; /* release waiters*/  
    }  
    else{  
        UNLOCK(bar_name.lock);  
        while (bar_name.flag != local_sense) {};  
    }  
}
```

Summary

- Locks and barriers for synchronization
- Low-level primitives in hardware
- Sophisticated synchronization algorithms in software